

Post-Internal Roadmap

What changes after the internal round — SIH26017 Land Acquisition Delay Prediction

The internal-round build is a deliberate beta — scoped for a 7-day timeline, not the final product. This roadmap is worth having ready for Q&A;: it shows judges we know exactly what we simplified, and why, and that we have a credible plan to close each gap.

Current (Beta) vs Planned (Next Phase)

Area	Current (Beta)	Planned (Next Phase)
Database	SQLite — zero setup, fast to build, fine for a single-user demo	PostgreSQL — supports concurrent multi-user access and stronger data governance, appropriate for sensitive government records
Community sentiment	Synthetic proxy score, correlated with consent/SIA status	Real NLP sentiment analysis run on actual SIA reports and public consultation records
Data source	Synthetic dataset, calibrated to real published patterns (NJDG, DoLR, DILRMP)	Direct integration with real government data — LACRRIS, NJDG, DILRMP-MIS — via API where available
Geospatial view	Not included in this phase; state/district analytics used instead	GIS-enabled map showing high-risk projects geographically, as named in the official ask
Alerts	Static mock panel only	Real automated notifications (email/SMS) to project managers and administrators for newly high-risk projects
Access control	None — open demo access	Secure, role-based login for different stakeholders (officer, policymaker, admin), with audit trails
Model learning	Trained once on the synthetic dataset	Continuous learning pipeline — model retrains as new project data becomes available, per the official ask
Deployment	Local / demo-only	Hosted deployment with proper environment separation (staging/production)

Why frame it this way to judges

- Every simplification in the beta has a stated, deliberate reason — timeline, not lack of understanding.
- Each “next phase” item maps directly to something in the official problem statement’s expected-solution list — showing the roadmap isn’t invented, it’s closing known gaps.
- This is a stronger answer to “why doesn’t your system do X yet” than pretending the beta is complete.

Suggested Q&A; line

“For this round, we prioritized building and validating the core predictive model and explainability logic within our timeline. The database, real sentiment analysis, GIS layer, and production security are all planned next steps — we scoped this deliberately rather than building a shallow version of everything.”